

Project 1: Deep Learning-Image Classification with CNN

Group 3

Presented by

Mauricio


Nicolas

How the partnership will work



Introduction

The main goal of this project is to develop an image classification model capable of correctly identifying 10 different animal species.



Data set

The Animals-10 dataset used in this project contains approximately 28,000 images distributed across 10 different animal classes: dog, cat, horse, elephant, butterfly, chicken, cow, sheep, spider, and squirrel.

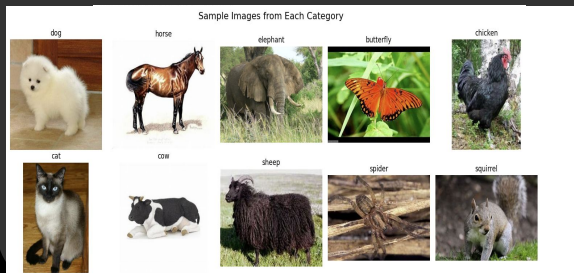
```
import kagglehub

# Download latest version
path = kagglehub.dataset_download("alessiocrorad  
99/animals10")

print("Path to dataset files:", path)
```

Visualization / Analysis

1



Preprocessing

2

CONFIGURATION PARAMETERS:

TARGET_SIZE: (224, 224)

BATCH_SIZE: 64

VALIDATION_SPLIT: 0.2

SEED: 42

EPOCHS: 25

LEARNING_RATE: 0.001

MODEL_PATH: /content/drive/MyDrive/animals_project/models/best_model.h5

Augmentation

3

```
train_datagen = ImageDataGenerator(  
    preprocessing_function=preprocess_input,  
    rotation_range=15,  
    width_shift_range=0.15,  
    height_shift_range=0.15,  
    shear_range=0.15,  
    zoom_range=0.15,  
    brightness_range=[0.8, 1.2],  
    horizontal_flip=True,  
    fill_mode='nearest',  
    validation_split=VALIDATION_SPLIT  
)
```

Validation Split

4

- 80% for training
- 20% for testing



MODELS AND TESTING

Architecture: 3 Conv+Pooling blocks → Flatten → Dense(256)+Dropout → Output
Filters: 32 → 64 → 1180 (aggressive expansion)
Activation: ReLU (hidden), Softmax (output)
Regularization: 50% Dropout
Optimization: Adam (lr=0.0005) + Categorical Crossentropy
I/O: 180x180x3 input → 10-class output
Parameters: ~8.5M
Performance: ~60% Accuracy

60%



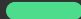
Architecture: 3 Conv+Pooling blocks → Flatten → Dense(256)+Dropout → Output
Filters: 32 → 64 → 1180 (aggressive expansion)
Activation: ReLU (hidden), Softmax (output)
Regularization: 50% Dropout
Optimization: Adam (lr=0.0002) + Categorical Crossentropy
I/O: 224x224x3 input → 10-class output
Parameters: ~8.5M
Performance: ~80% Accuracy

80%

Architecture: 4 Conv+Pooling blocks → Flatten → Dense+Dropout → Output
Filters: 32 → 64 → 128 → 128 (progressive increase)
Activation: ReLU (hidden), Softmax (output)
Regularization: 50% Dropout
Optimization: Adam + Categorical Crossentropy
I/O: 224x224x3 input → 10-class output

75%

Transfer Learning

-  EfficientNetB0
-  Resnet50
-  MobileNetV2

Fine - Tuning

```
# Unfreeze top layers for fine-tuning (keep early layers frozen)
FINE_TUNE_AT = 100 # For GPU L4, this is a safe unfreeze depth
for layer in model_tl.layers[:FINE_TUNE_AT]:
    layer.trainable = False
for layer in model_tl.layers[FINE_TUNE_AT:]:
    layer.trainable = True
```

98%

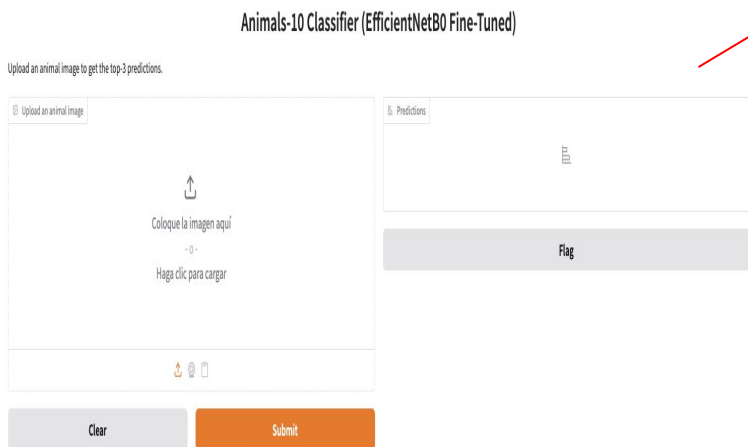
VALIDATION RESULTS:

- Validation Accuracy: 0.9784
- Validation Loss: 0.0722

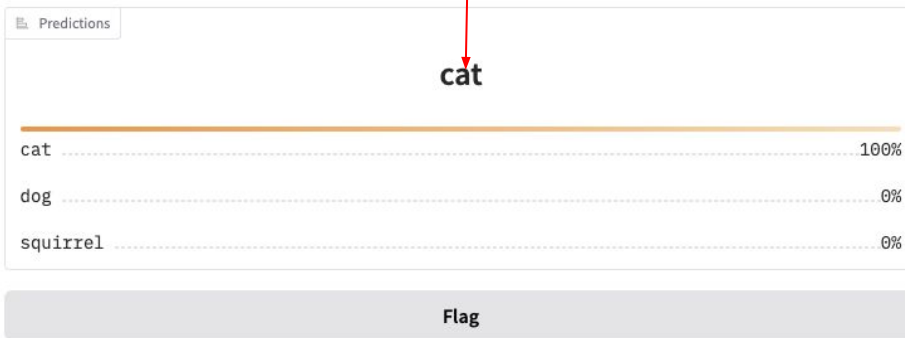
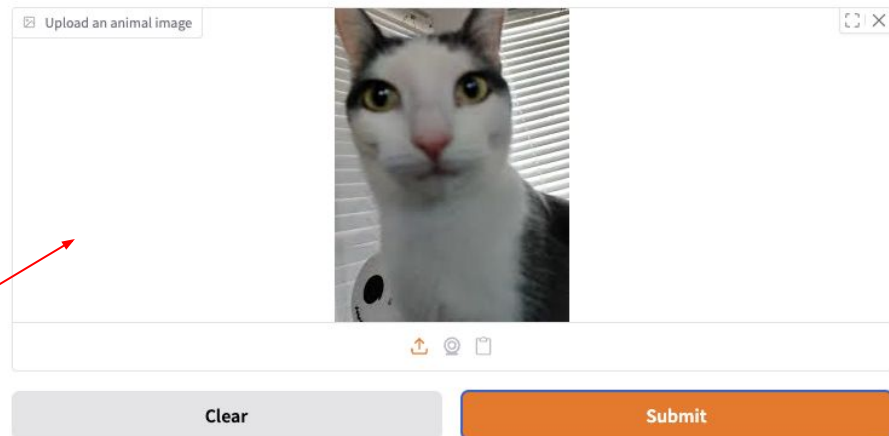
Let's work together

Deployment

- We use Gradio for the deployment.



Upload an animal image to get the top-3 predictions.



RESULTS

VALIDATION RESULTS:

- Validation Accuracy: 0.9784
- Validation Loss: 0.0722

CLASSIFICATION REPORT:

	precision	recall	f1-score	support
cane	0.99	0.98	0.98	972
cavallo	0.97	0.98	0.97	524
elefante	0.95	0.99	0.97	289
farfalla	0.98	0.99	0.98	422
gallina	0.99	0.99	0.99	619
gatto	0.98	0.96	0.97	333
mucca	0.93	0.97	0.95	373
pecora	0.97	0.94	0.95	364
ragno	0.99	0.99	0.99	964
scoiattolo	0.98	0.97	0.98	372
accuracy			0.98	5232
macro avg	0.97	0.98	0.97	5232
weighted avg	0.98	0.98	0.98	5232

